

Graph Neural Networks

Natalia Dubljevic
Biomedical Engineering
Schulich School of Engineering

W2024

Outline


- Learning Goals
- Graph neural networks
 - What is graph-structured data
 - What sorts of problems can we solve using graph neural networks
 - How are graph neural networks designed
- Summary
- Tutorial

Learning Goals

- Understand graph structured data and its differences from data we've seen so far
- What are some applications of GNNs
- Some of the core concepts of GNN implementation such as
 - Message passing
 - K-hop neighbourhood

What is a graph anyways?

- Graphs have **nodes** and **edges**
- Edges represent the relationship between any two nodes



Graph structured data


Graphs can be used to represent data such as:

Social networks



[University of Kentucky]

Molecules




[flaticon.com]

Citation Maps



[connectedpapers.com]

Tissues



Transit

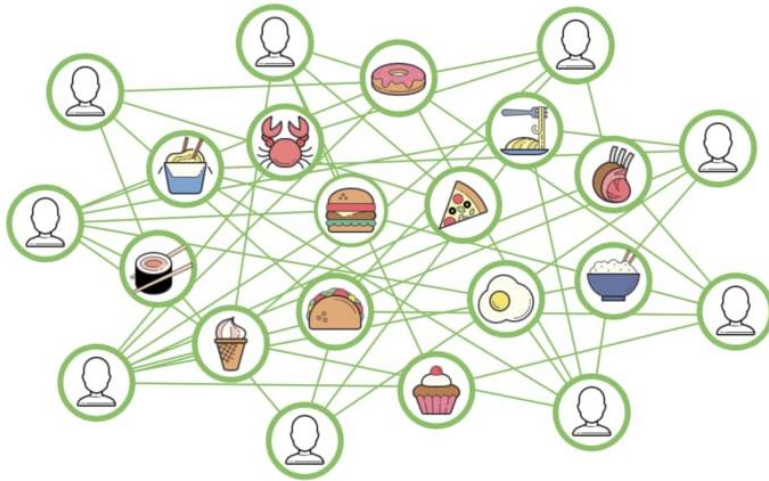


[translink.ca]

Graph structured data

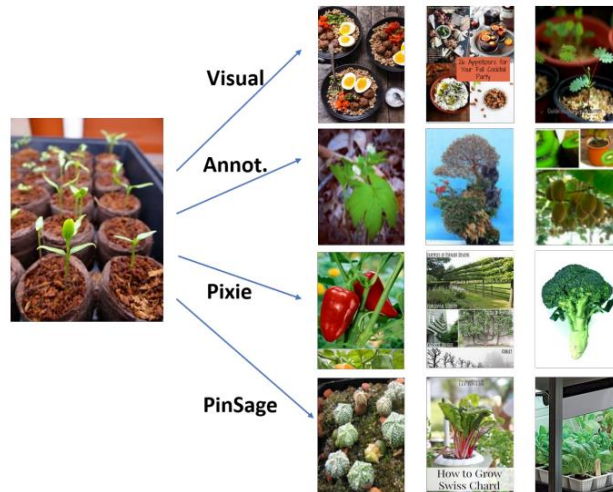
GNNs aren't just for niche applications-- they're trendy in industry too!

Uber Eats




[uber.com]

Pinterest



[Ying et al. 2018]

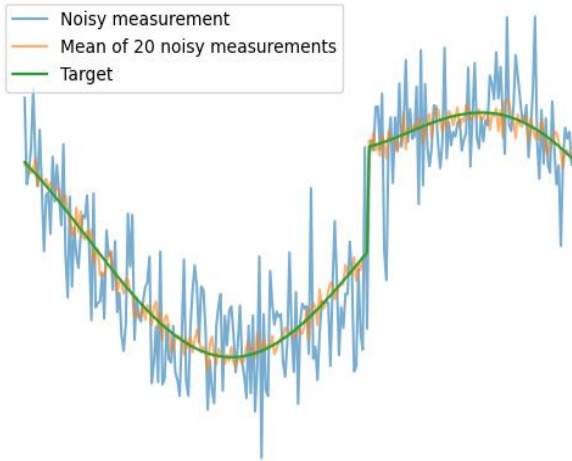
Google maps



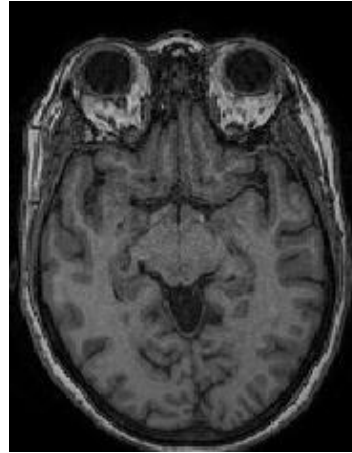
[deepmind.google.com]

What we've seen so far

Signals



Images



Tabular

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife


Text

Love looks not with the eyes, but
with the mind; And therefore is
wing'd Cupid painted blind. Nor
hath love's mind of any judgment
taste; Wings and no eyes figure
unheedy haste: And therefore is
love said to be a child, Because in
choice he is so oft beguil'd.

- All of these can be represented as graphs!
- Graphs are one of the most general data structures

A bit more on graphs


- We can represent graphs as an adjacency matrix



	A	B	C	D	E
A	0	1	1	1	0
B	1	0	0	0	1
C	1	0	0	0	0
D	1	0	0	0	1
E	0	1	0	1	1


A bit more on graphs

- We can represent graphs as an adjacency matrix
- Edges can be weighted and/or directional



A bit more on graphs


- We can represent graphs as an adjacency matrix
- Edges can be weighted and/or directional



	A	B	C	D	E
A	0	1	1	0	0
B	1	0	0	0	1
C	1	0	0	0	0
D	1	0	0	0	1
E	0	1	0	1	1

A bit more on graphs

- We can represent graphs as an adjacency matrix
- Edges can be weighted and/or directional
- Can have features related to nodes, edges, or graph as a whole
 - Node features most common



	A	B	C	D	E
A	0	1	1	1	0
B	1	0	0	0	1
C	1	0	0	0	0
D	1	0	0	0	1
E	0	1	0	1	1

Working with sparse tensors

- Since they are often sparse, adjacency matrices are commonly represented as a list of edges

	A	B	C	D	E
A	0	1	1	1	0
B	1	0	0	0	1
C	1	0	0	0	0
D	1	0	0	0	1
E	0	1	0	1	1

Working with sparse tensors

- Since they are often sparse, adjacency matrices are commonly represented as a list of edges

	0	1	2	3	4
0	0	1	1	1	0
1	1	0	0	0	1
2	1	0	0	0	0
3	1	0	0	0	1
4	0	1	0	1	1

Working with sparse tensors

- Since they are often sparse, adjacency matrices are commonly represented as a list of edges

	0	1	2	3	4
0	0	1	1	1	0
1	1	0	0	0	1
2	1	0	0	0	0
3	1	0	0	0	1
4	0	1	0	1	1

$(0, 1)$

Working with sparse tensors

- Since they are often sparse, adjacency matrices are commonly represented as a list of edges

	0	1	2	3	4
0	0	1	1	1	0
1	1	0	0	0	1
2	1	0	0	0	0
3	1	0	0	0	1
4	0	1	0	1	1

$$\begin{pmatrix} (0, 1) \\ (0, 2) \end{pmatrix}$$

Working with sparse tensors

- Since they are often sparse, adjacency matrices are commonly represented as a list of edges

	0	1	2	3	4
0	0	1	1	1	0
1	1	0	0	0	1
2	1	0	0	0	0
3	1	0	0	0	1
4	0	1	0	1	1

$$\begin{pmatrix} (0, 1) \\ (0, 2) \\ (0, 3) \end{pmatrix}$$

Working with sparse tensors

- Since they are often sparse, adjacency matrices are commonly represented as a list of edges

	0	1	2	3	4
0	0	1	1	1	0
1	1	0	0	0	1
2	1	0	0	0	0
3	1	0	0	0	1
4	0	1	0	1	1

$$\begin{pmatrix} (0, 1) \\ (0, 2) \\ (0, 3) \\ (1, 0) \end{pmatrix}$$

Working with sparse tensors


- Since they are often sparse, adjacency matrices are commonly represented as a list of edges

	0	1	2	3	4
0	0	1	1	1	0
1	1	0	0	0	1
2	1	0	0	0	0
3	1	0	0	0	1
4	0	1	0	1	1


$$\begin{pmatrix} (0, 1) \\ (0, 2) \\ (0, 3) \\ (1, 0) \\ (1, 4) \\ \vdots \end{pmatrix}$$

Tasks


Node Classification




Edge prediction



Clustering




Graph Classification



Prelude to message passing


- How do we create models that can accommodate graph-structured data?
- Naïve approach: Feed a node and its features into an FCNN



- **Question:** Are there any issues with this approach?

Prelude to message passing

- Unlike standard supervised learning, we do not have independent points!
 - Our nodes are connected to each other in unique ways




is not the
same as




- We want to develop models that can exploit the relationships between the nodes
- We'll introduce '**message passing**' which allows for information to be propagated throughout a neural network

Message passing




Message passing

1. **Message:** Nodes send messages to their neighbors




Message passing

1. **Message:** Nodes send messages to their neighbors
2. **Aggregate:** Messages are aggregated in a permutation-invariant way




Message passing

1. **Message:** Nodes send messages to their neighbors
2. **Aggregate:** Messages are aggregated in a permutation-invariant way
3. **Update:** Node embeddings are updated




Message passing

- We just focused on updating the pink node, but in practice, all nodes are updated at the same time!




Message passing

- We just focused on updating the pink node, but in practice, all nodes are updated at the same time!




Message passing

- We just focused on updating the pink node, but in practice, all nodes are updated at the same time!




Message passing

- We just focused on updating the pink node, but in practice, all nodes are updated at the same time!





k-hop neighborhood

- The number of layers we have corresponds to the 'k-hop' neighborhood





k-hop neighborhood

- The number of layers we have corresponds to the 'k-hop' neighborhood





k-hop neighborhood

- The number of layers we have corresponds to the 'k-hop' neighborhood





k-hop neighborhood

- The number of layers we have corresponds to the 'k-hop' neighborhood



k-hop neighborhood


- The number of layers we have corresponds to the 'k-hop' neighborhood



k-hop neighborhood


- The more hops, the more distant information each node gets
- More hops is good for global context
- ...but too many hops will wash out local structural information

2-hop subgraph
1-hop subgraph




Trainable parameters

- Can think of it as two places we can have trainable parameters:




Trainable parameters

- Can think of it as two places we can have trainable parameters:
 - When passing a message



Trainable parameters


- Can think of it as two places we can have trainable parameters:
 - When passing a message
 - When updating an embedding



Trainable parameters

- Can think of it as two places we can have trainable parameters:
 - When passing a message
 - When updating an embedding

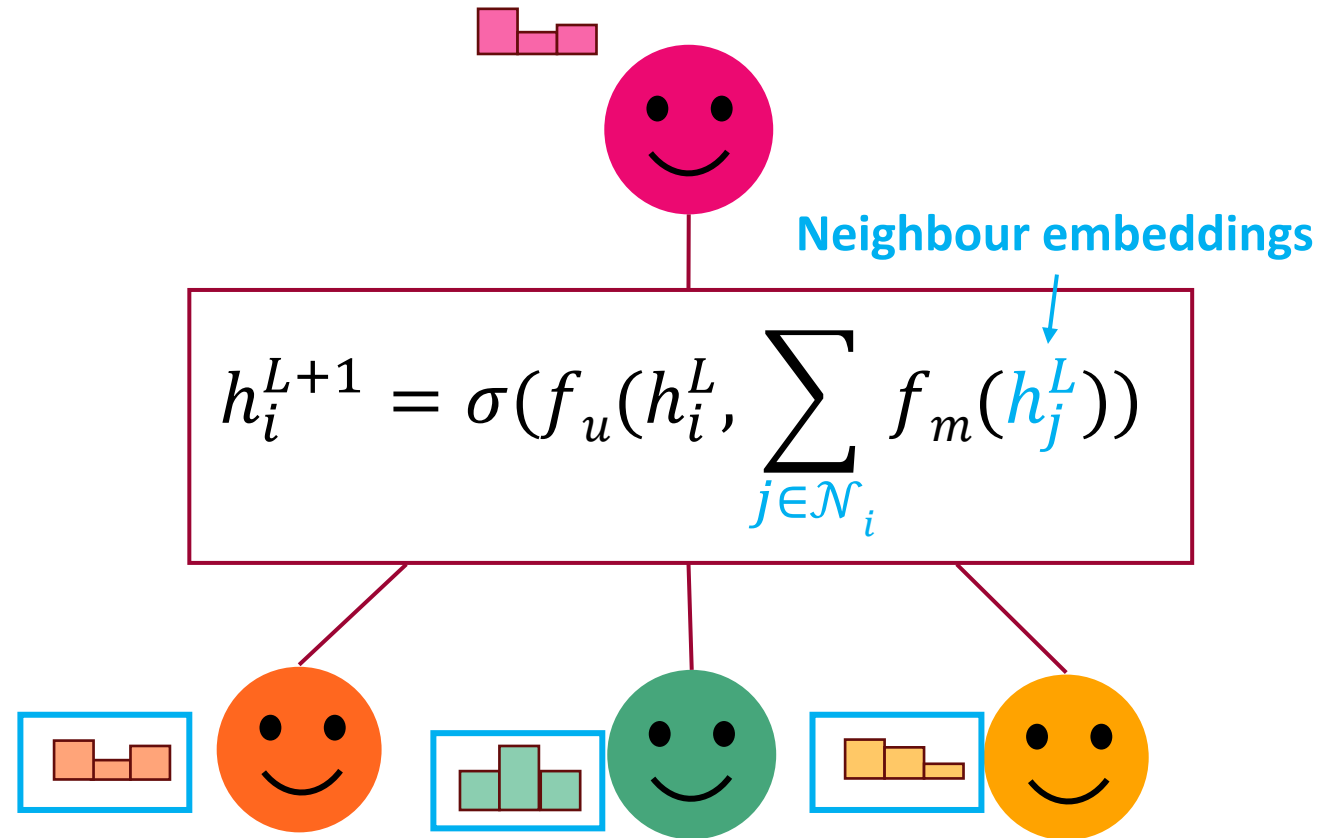
* Note: The trainable parameters are shared across ALL nodes!



Trainable parameters

- Can think of it as two places we can have trainable parameters:
 - When passing a message
 - When updating an embedding

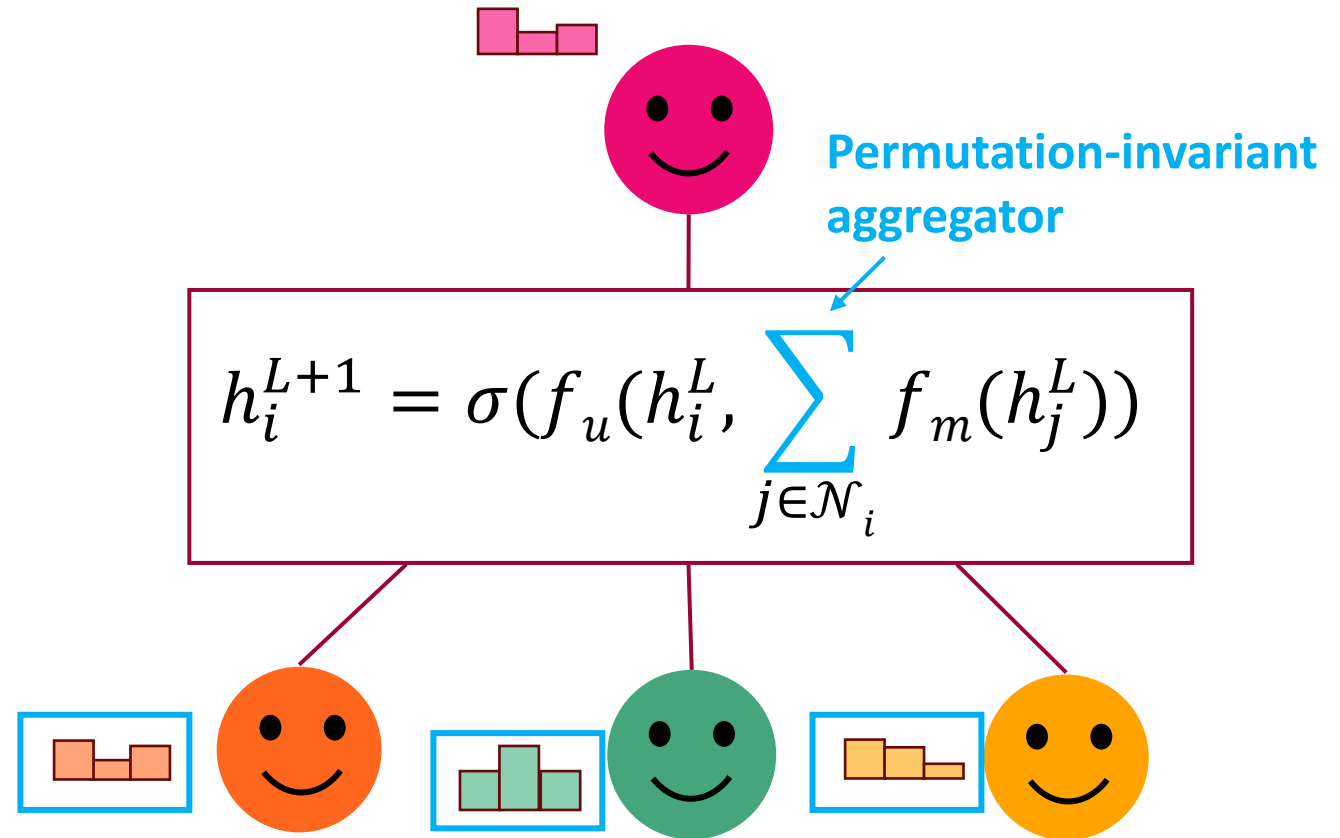
* Note: The trainable parameters are shared across ALL nodes!



Trainable parameters

- Can think of it as two places we can have trainable parameters:
 - When passing a message
 - When updating an embedding

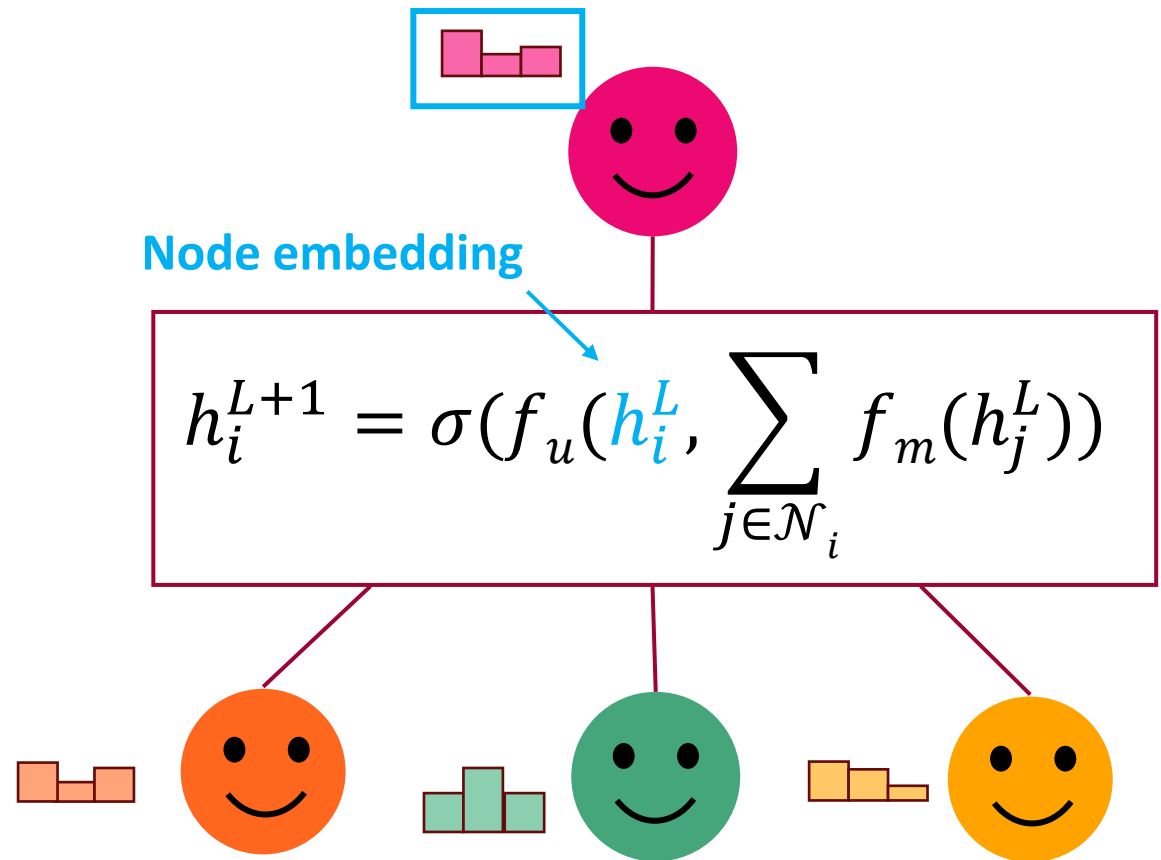
* Note: The trainable parameters are shared across ALL nodes!



Trainable parameters

- Can think of it as two places we can have trainable parameters:
 - When passing a message
 - When updating an embedding

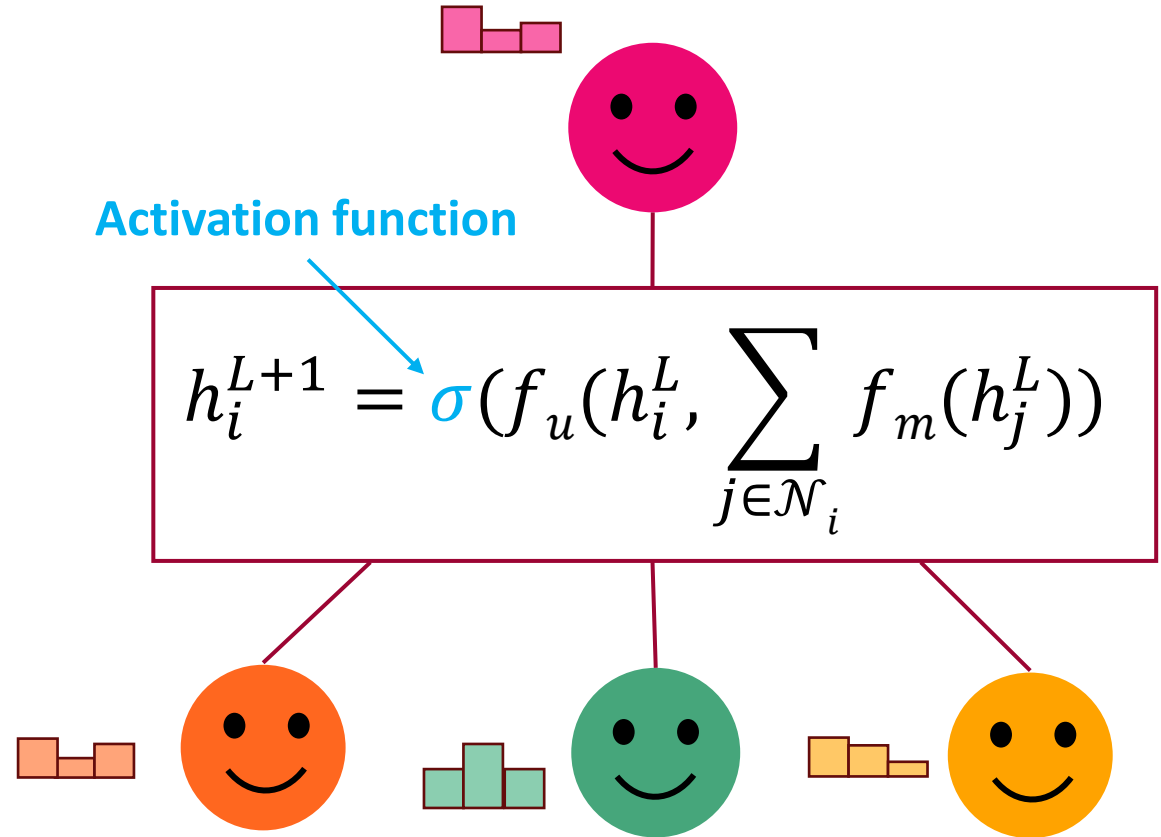
* Note: The trainable parameters are shared across ALL nodes!



Trainable parameters

- Can think of it as two places we can have trainable parameters:
 - When passing a message
 - When updating an embedding

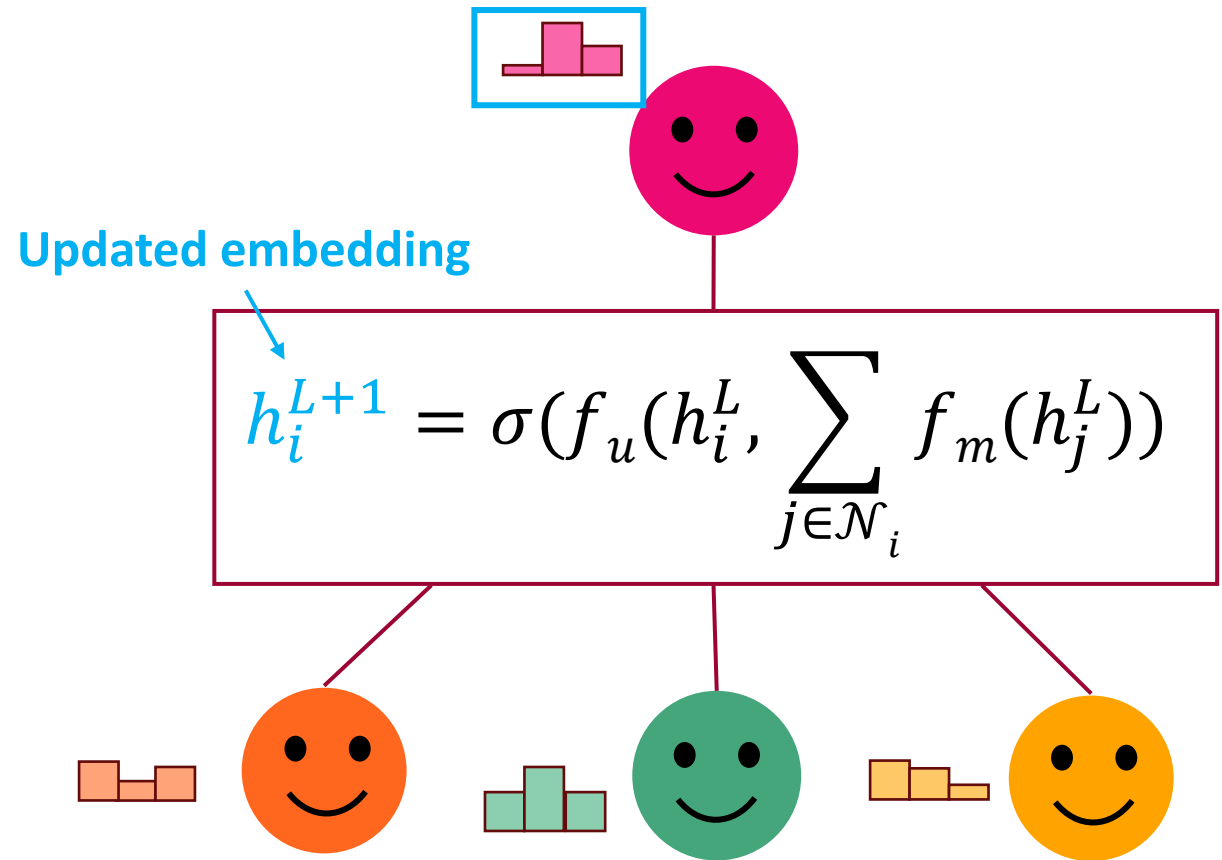
* Note: The trainable parameters are shared across ALL nodes!



Trainable parameters


- Can think of it as two places we can have trainable parameters:
 - When passing a message
 - When updating an embedding

* Note: The trainable parameters are shared across ALL nodes!




GNN layer example: Graph Convolutional Network

- Graph convolutional network (GCN) learns a simple matrix of weights
 - Can think of this matrix as a 'convolutional kernel' we use at each node instead of pixel



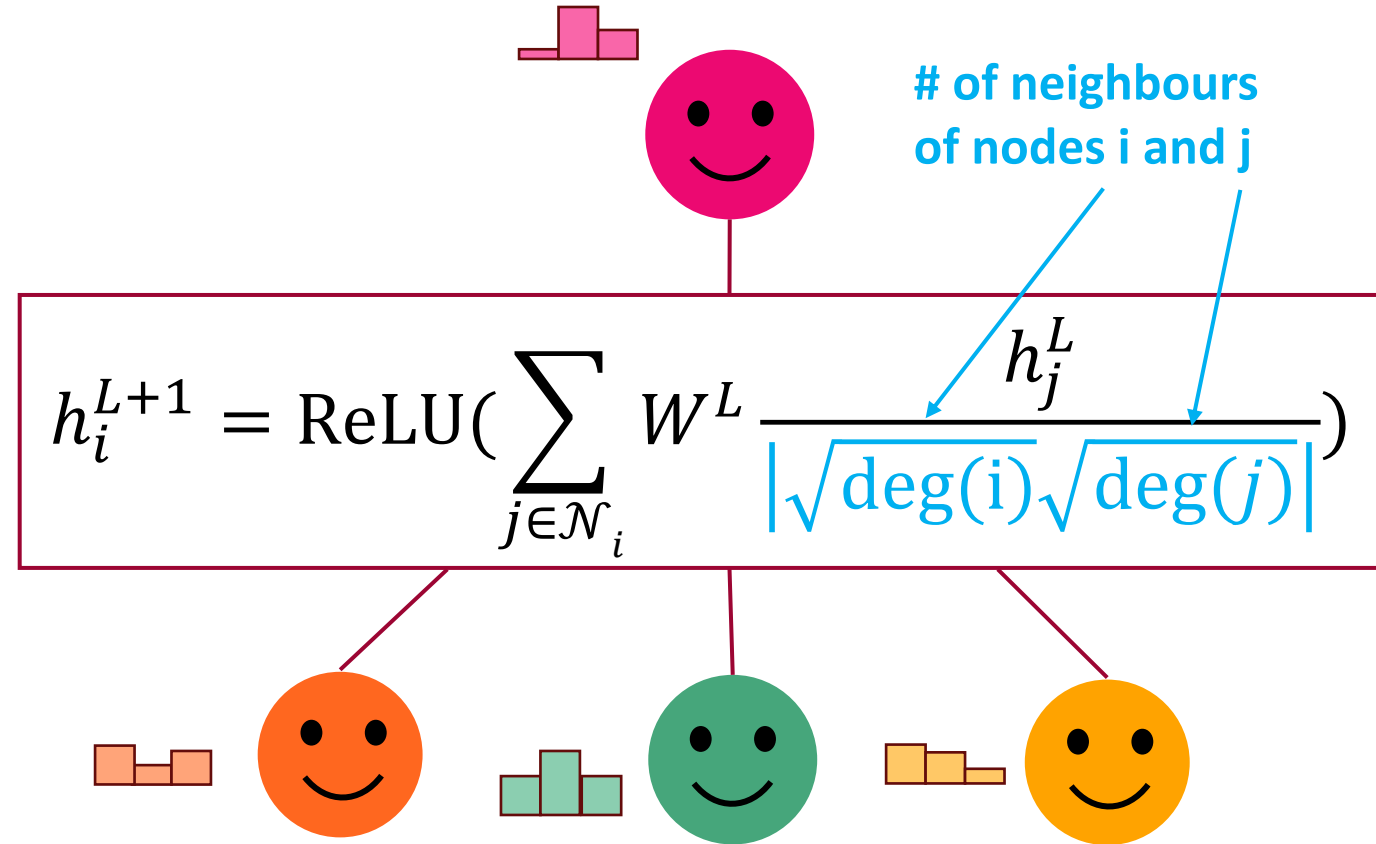
GNN layer example: Graph Convolutional Network

- Graph convolutional network (GCN) learns a simple matrix of weights
 - Can think of this matrix as a 'convolutional kernel' we use at each node instead of pixel



GNN layer example: Graph Convolutional Network

- Graph convolutional network (GCN) learns a simple matrix of weights
 - Can think of this matrix as a 'convolutional kernel' we use at each node instead of pixel



The diagram illustrates the GCN layer operation. At the top, a pink node is connected to three nodes below: orange, green, and yellow. Each node has a small bar chart next to it representing its features. The pink node's bar chart has three bars of increasing height. The orange node's bar chart has three bars of decreasing height. The green node's bar chart has three bars of increasing height. The yellow node's bar chart has three bars of decreasing height. The equation for the next layer's node feature is shown in a box:


$$h_i^{L+1} = \text{ReLU}\left(\sum_{j \in \mathcal{N}_i} W^L \frac{h_j^L}{|\sqrt{\deg(i)}\sqrt{\deg(j)}|}\right)$$

Annotations in the diagram include:

- A blue arrow pointing from the text "# of neighbours of nodes i and j" to the denominator of the equation.
- A blue arrow pointing from the text h_j^L to the numerator of the equation.

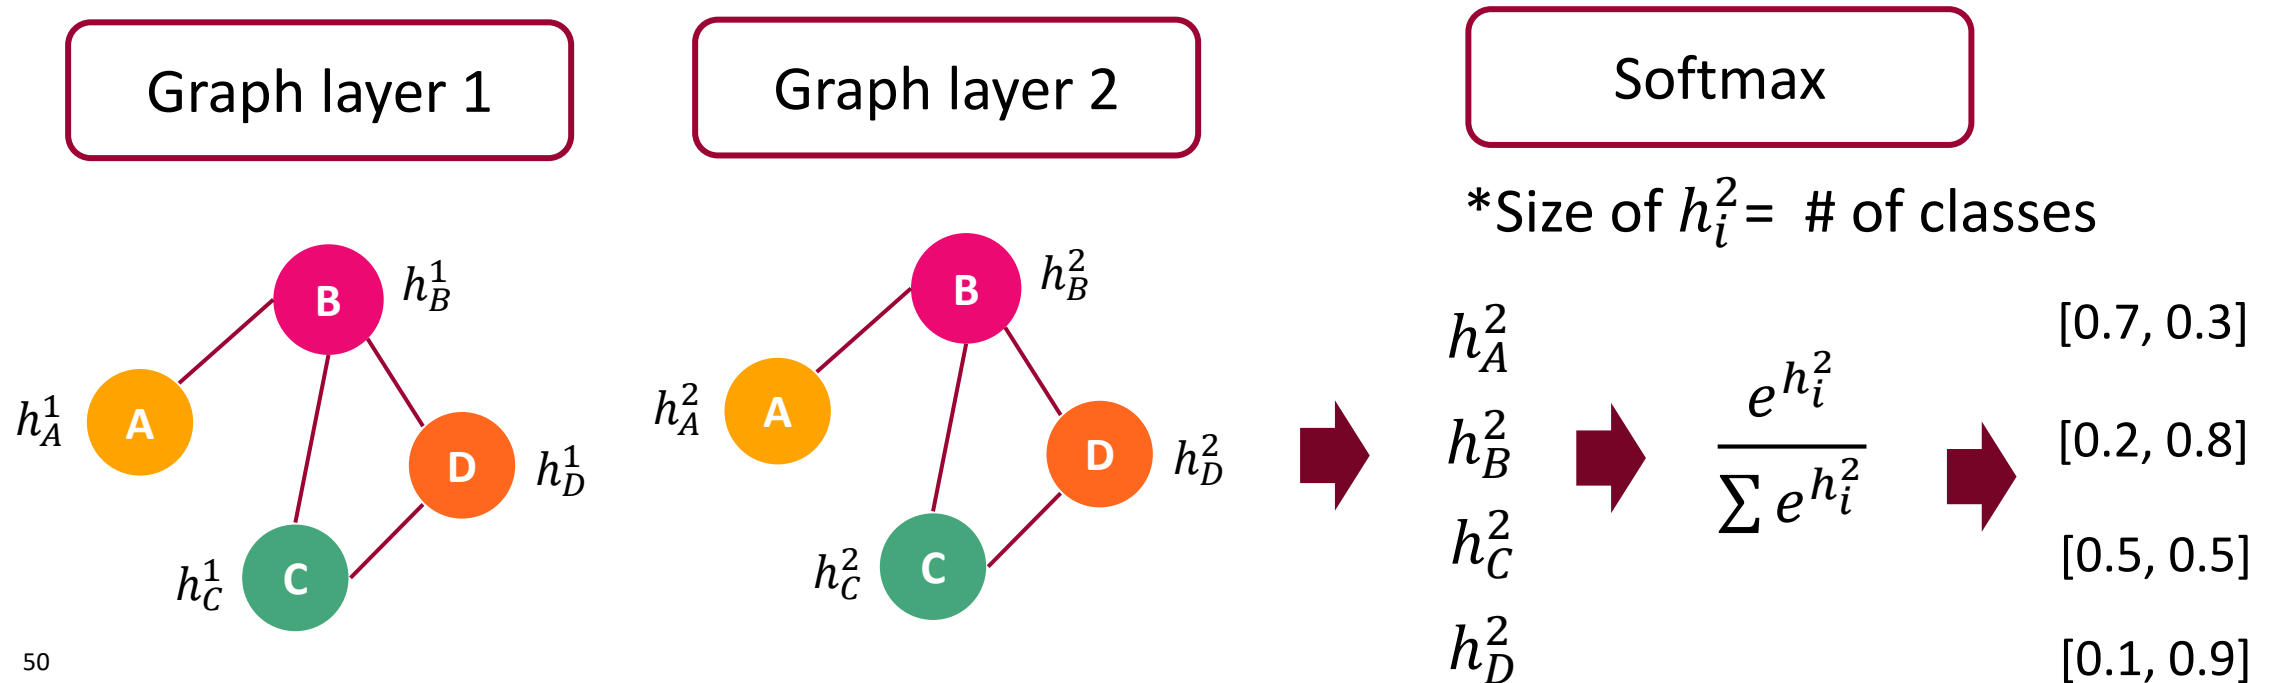
GNN layer example: Graph Convolutional Network

- Graph convolutional network (GCN) learns a simple matrix of weights
 - Can think of this matrix as a 'convolutional kernel' we use at each node instead of pixel
- GCN assumes self-edges!
 - So, each node is its own neighbor 🤔



Putting it all together

- We *could* just learn node/edge/graph embeddings e.g. node2vec
- To tailor network for a specific task, we add an appropriate head
- E.g. Node classification:




Summary

- Graph structured data is one of the most general types of data that encompasses a broad range of applications
- Graph neural networks can be used to solve a variety of tasks
- Since graph data is inherently different from other types of data, we need a unique framework to handle it
- Message passing allows for all the information in a graph to be effectively harnessed


Next class:

- We will learn how to work with graphs and build models using pytorch-geometric!

Cora dataset (Citations)



TU MUTAG Dataset (Molecules)



References

- J. Leskovec. [CS224W: Machine Learning with Graphs](#). ** Great option if you want a whole COURSE on graph-based machine learning
- W. L. Hamilton. Graph Representation Learning. *Synthesis Lectures on Artificial Intelligence and Machine Learning*, 14(3), pages 1-159. 2020.
- Sanchez-Lengeling, et al., "A Gentle Introduction to Graph Neural Networks", Distill, 2021.
- T. N. Kipf et al., Semi-Supervised Classification with graph Convolutional Networks, ICLR 2017.
- K. Xu et al., How Powerful are Graph Neural Networks? ICLR, 2019.
- W. L. Hamilton et al., Inductive representation Learning on Large Graphs. NIPS, 2017.
- J. Zhou et al., Graph neural networks: A review of methods and applications. AI Open 1. Pages 57-81. 2020.
- Daigavane, et al., "Understanding Convolutions on Graphs", Distill, 2021.
- P. Lippe. [Tutorial 7: Graph Neural Networks — UvA DL Notebooks v1.2 documentation \(uvadlc-notebooks.readthedocs.io\)](#). 2022.
- M. N. Bernstein. [Graph convolutional neural networks - Matthew N. Bernstein \(mbernste.github.io\)](#). 2023.
- Z. Wu et al., A Comprehensive Survey on Graph Neural Networks. IEEE Trans. On Neural networks and Learning Systems. 32(1). 2021.
- R. Anand. [Math Behind Graph Neural Networks - Rishabh Anand \(rish-16.github.io\)](#). 2022.
- D. Grattarola. [A practical introduction to GNNs - Part 2 – Daniele Grattarola](#). 2021.
- E. Benjaminson. [Understanding Message Passing in GNNs – Emma Benjaminson – Data Scientist \(sassafra13.github.io\)](#). 2022.
- T. Masui. [Graph Neural Networks with PyG on Node Classification, Link Prediction, and Anomaly Detection | by Tomonori Masui | Towards Data Science](#). 2022.

Thank you!

